

Gabriel Ryan

gabe@cs.columbia.edu | www.cs.columbia.edu/~gabe

Research Focus:

My research interests are at the intersection of Machine Learning, Programming Languages, and Software Engineering/Security. My recent work focuses on using LLMs and probabilistic methods for program analysis and testing.

Education:

Ph.D. Columbia University

Anticipated December 2023

- Ph.D. Candidate in Computer Science advised by Professor Suman Jana.

M.S. Columbia University

December 2017

- M.S. in Computer Science
- Conducted research in modeling and simulating user behavior for Intrusion Detection Lab.
- Selected Coursework: Security, Deep Learning, Adv. Deep Learning, Data Analytics, Data Exploration Systems, Machine Learning, Natural Language Processing, Operating Systems, Compilers, High Perf. Computing

B.S. Swarthmore College

June 2013

- B.S. Engineering, B.A. Computer Science
- Selected Coursework: Adv. Algorithms, Theory of Computation, Artificial Intelligence, Probabilistic Robotics, Computer Systems, Computer Languages, Applied Statistics, Optimization, Linear Systems Analysis

Research Experience:

Doctoral Researcher, Columbia Security Lab

Sep 2018 - Present

- Developing new approaches to linux kernel syzbot crash report explanation & bug repair with LLMs.
- Developed novel linux kernel race prediction analysis with probabilistic and spectral learning. Published in Oakland S&P 2023.
- Developed *Proximal Gradient Analysis* (PGA), a new method for dataflow analysis that uses nonsmooth optimization methods. Implemented in LLVM. Published in USENIX Security 2021.
- Designed the *Continuous Logic Network* (CLN), a neural architecture for logical learning and reasoning, and applied it to learning program invariants for verification. Implemented library for constructing CLNs in Pytorch. Published in ICLR 2020 and PLDI 2020.

Research Intern, AWS A.I. Labs

Jun 2023 - Aug 2023

- Research internship with AWS CodeWhisperer research group advised by visiting Prof. Baishahki Ray.
- Developed novel approach to regression testing with LLMs using static analysis to prompt the model to reason symbolically about program execution paths.
- Achieved improvements of $2\times$ coverage and $3\times$ correct test generations over baseline LLM generation approaches when evaluated on CodeGen2 and OpenAI GPT-3.5 and GPT-4 models. Submitted conference paper currently under review.

Research Intern, Microsoft Research

Jun 2021 - Aug 2021

- Research internship with MSR RiSE group advised by Sr. Pr. Scientists Todd Mytkowitz and Shuvendu Lahiri.
- Developed *TOGA: A Neural Method for Test Oracle Generation* using Transformers (CodeBERT) and a specialized grammar to generate unit tests and demonstrated a 170% improvement over prior work in bug detection.
- Published work in ICSE 2022, awarded ACM SigSoft Distinguished Paper Award.

MS Graduate Research Assistant, Columbia IDS Lab

Fall 2016 - Fall 2017

- Developed novel method for modeling user behavior with neural networks advised by Prof. Salvatore Stolfo.
- Designed architecture combining recurrent neural networks with multi task prediction to model events and times in a real-time sequence. Published in IEEE S&P Workshops 2018.

Research Intern, Woods Hole Oceanographic Institute

January 2012

- Conducted frequency analysis of velocity data and developed burst sampling system that improved the signal to noise ratio by a factor of 3 while minimizing computation and radio usage.
- Published work in proceedings of Oceans 2012 conference.

Professional Experience:

Software Engineer, Allure Security Technology

Sep 2015 - August 2018

- Developed web application for managing large scale deployments of User Behavior sensors and data loss detection using Java Spring Framework with AngularJS front end and Hibernate with Postgres and MongoDB databases. Designed and implement REST API using Swagger specification and OAuth2 authentication.
- Developed midstream network based document interception and tracking system using Squid Proxy and ICAP.

Robotics Software Consultant, 3DDataLtd

May 2015 - Aug 2015

- Developed a framework for 3d Simultaneous Localization and Mapping using the Lidar Odometry and Mapping algorithm. Modified algorithm to integrate IMU Sensor data using an Extended Kalman Filter.

Radar Software Engineer, Raytheon

Sep 2013 - April 2015

- Analyzed data from radar testing and provide software support for automated calibration, satellite tracking, antenna diagnostics, and maintenance prioritization software in Ada and C++.
- Earned team achievement award for completing new diagnostic capabilities ahead of schedule.

Publications:

[Oakland S&P 2023] *Precise Detection of Kernel Data Races with Probabilistic Lockset Analysis*. G. Ryan, A. Shah, D. She, S. Jana. https://cs.columbia.edu/~gabe/files/oakland2023_pla.pdf

[ICSE 2022] *TOGA: A Neural Method for Test Oracle Generation*. E. Dinella*, G. Ryan*, T. Mytkowitz, S. Lahiri. <https://arxiv.org/pdf/2109.09262.pdf> (**Distinguished Paper Award**)

[OSDI 2021] *DistAI: Data-Driven Automated Invariant Learning for Distributed Protocols*. J. Yao, R. Tao, R. Gu, J. Nieh, S. Jana, G. Ryan. <https://www.usenix.org/system/files/osdi21-yao.pdf> (**Best Paper Award**)

[USENIX Security 2021] *Fine Grained Dataflow Tracking with Proximal Gradients*. G. Ryan, A. Shah, D. She, K. Bhat, and S. Jana. <https://arxiv.org/abs/1909.03461>

[PLDI 2020] *Learning Nonlinear Loop Invariants with Gated Continuous Logic Networks*. J. Yao*, G. Ryan*, J. Wong*, S. Jana, and R. Gu. <https://arxiv.org/abs/2003.07959>

[ICLR 2020] *CLN2INV: Learning Loop Invariants with Continuous Logic Networks*. G. Ryan*, J. Wong*, Y. Jianan*, R. Gu, and S. Jana. <https://arxiv.org/abs/1909.11542>

[InfoVis 2018] *At a Glance: Pixel Approximate Entropy as a Measure of Line Chart Complexity*. G. Ryan, A. Mosca, R. Chang, and E. Wu. <https://arxiv.org/abs/1811.03180>

[Oakland S&P Workshops 2018] *Simulated User Bots: Real Time Testing of Insider Threat Detection Systems*. P. Dutta, G. Ryan, A. Zieba, and S. Stolfo. <https://ieeexplore.ieee.org/document/8424654>

[Oceans 2012] *Oversampling MAVS for reduction of vortex-shedding velocity-sensing noise*. A. Williams, G. Ryan, and T. Thwaites. <https://ieeexplore.ieee.org/document/6404777>

Awards:

National Defense Science and Engineering Graduate Fellowship (NDSEG). Awarded NDSEG Fellowship for proposal "Proximal Gradient Analysis for Vulnerability Detection and Defense." 2019.

NSF Graduate Research Fellowship Honorable Mention. Accorded honorable mention for proposal "Modeling and Simulating Adversarial User Behavior with Sequential VAEs." 2018.

Academic Service:

AAAI 2023. Program Committee.

TOSEM 2022. Reviewer.

AAAI 2022. Reviewer.

Teaching:

Teaching Assistant. Intrusion Detection Systems. Fall 2022.

Lecturer. Continuous Logic Networks. Spring 2021.

Teaching Assistant. Intrusion Detection Systems. Fall 2020.

Technical Summary:

Languages: Proficient in Python, C/C++, Java, Javascript, and SQL.

Skills: Linux, Cloud Compute, Transformers, Pytorch, Tensorflow, LLVM, Pandas, Z3